

Final Project Submission

Please fill out:

- Student name:
- Student pace: self paced / part time / full time:
- Scheduled project review date/time:
- Instructor name:
- Blog post URL:
- Video of 5-min Non-Technical Presentation:

TABLE OF CONTENTS

Click to jump to matching Markdown Header.

- [Introduction](#)
 - [OBTAIN](#)
 - [SCRUB](#)
 - [EXPLORE](#)
 - [MODEL](#)
 - [INTERPRET](#)
 - [Conclusions/Recommendations](#)
-

INTRODUCTION

Explain the point of your project and what question you are trying to answer with your modeling.

Business Problem

Summary of the business problem you are trying to solve, and the data questions that you plan to answer to solve them.

Questions to consider:

- What are the business's pain points related to this project?
- How did you pick the data analysis question(s) that you did?
- Why are these questions important from a business perspective?

OBTAIN

Data Understanding

Describe the data being used for this project.

Questions to consider:

- Where did the data come from, and how do they relate to the data analysis questions?
 - What do the data represent? Who is in the sample and what variables are included?
 - What is the target variable?
 - What are the properties of the variables you intend to use?
-

Importing packages for importing data and exploratory visual analysis.

In [500]:

```
1 import pandas as pd
2 import seaborn as sns
3 # sns.set_theme(color_codes=True)
4 import matplotlib.pyplot as plt
5 import numpy as np
6
7 ## Preprocessing tools
8 from sklearn.model_selection import train_test_split, cross_val_predict, cross_val_score
9 from sklearn.preprocessing import MinMaxScaler, StandardScaler, OneHotEncoder
10 scaler = StandardScaler()
11 from sklearn.impute import SimpleImputer
12 from sklearn.pipeline import Pipeline
13 from sklearn.compose import ColumnTransformer
14 from imblearn.over_sampling import SMOTE, SMOTENC
15 from sklearn import metrics
16
17 ## Models & Utils
18 from sklearn.dummy import DummyClassifier
19 from sklearn.linear_model import LogisticRegression, LogisticRegressionCV
20 from sklearn.ensemble import RandomForestClassifier
21 from sklearn.svm import SVC
22
23 from time import time
24
25 import pandas as pd
26 import numpy as np
27 import matplotlib.pyplot as plt
28 import seaborn as sns
29 from sklearn.preprocessing import StandardScaler
30 from sklearn.model_selection import train_test_split
31 from sklearn.linear_model import LogisticRegression
32 from sklearn.metrics import classification_report
33 from sklearn.model_selection import cross_val_score
34 from xgboost import XGBClassifier
35 import warnings
36 warnings.filterwarnings(action='ignore')
37 from sklearn.neighbors import KNeighborsClassifier
38 from sklearn.tree import DecisionTreeClassifier
39 from sklearn.svm import SVC
40 from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier, BaggingClassifier
41 from sklearn.decomposition import PCA
42 from sklearn.model_selection import GridSearchCV
43
44 from sklearn.metrics import plot_confusion_matrix
45
46 from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

In [501]:

```

1  # Visualize the impact of a few key metrics on Hall of Fame inclusivity
2  def comparative_graph(s):
3      cat, num = 'inducted', s
4      fig, ax = plt.subplots(nrows=1, ncols=3, sharex=False, sharey=False, fi
5      fig.suptitle(s + ' vs Inducted', fontsize=20)
6
7      # Create a distribution graph to compare HOF inducted players against th
8      ax[0].title.set_text('density')
9      for i in df[cat].unique():
10         sns.distplot(df[df[cat]==i][num], hist=False, label=i, ax=ax[0])
11     ax[0].grid(True)
12
13     # Create a stacked bar graph containing 10 bins to help visualize the di
14     ax[1].title.set_text('bins')
15     breaks = np.quantile(df[num], q=np.linspace(0,1,11))
16     tmp = df.groupby([cat, pd.cut(df[num], breaks, duplicates='drop')]).size
17     tmp = tmp[df[cat].unique()]
18     tmp["tot"] = tmp.sum(axis=1)
19     for col in tmp.drop("tot", axis=1).columns:
20         tmp[col] = tmp[col] / tmp["tot"]
21     tmp.drop("tot", axis=1).plot(kind='bar', stacked=True, ax=ax[1], legend=
22
23     # Create a boxplot to compare HOF inducted players against those not ind
24     ax[2].title.set_text('outliers')
25     sns.boxplot(x=cat, y=num, data=df, ax=ax[2])
26     ax[2].grid(True)
27     plt.savefig(s)
28     plt.show();
29
30
31 # Create a new correlated dataframe with absolute value of a number,
32 def high_corr(df):
33     df_highcorr = df.corr().abs().stack().reset_index().sort_values(0, ascen
34     df_highcorr['Highly Correlated Pairs'] = list(zip(df_highcorr.level_0, d
35     df_highcorr.set_index(['Highly Correlated Pairs'], inplace = True)
36     df_highcorr.drop(columns=['level_1', 'level_0'], inplace = True)
37     df_highcorr.columns = ['Correlation']
38     df_highcorr.drop_duplicates(inplace=True)
39     return df_highcorr[(df_highcorr.Correlation>.7) & (df_highcorr.Correlati
40
41 # Create function used to find Precision, Recall, Accuracy, and F1 Scores.
42 def print_metrics(labels, preds):
43     print("Precision Score: {}".format(precision_score(labels, preds)))
44     print("Recall Score: {}".format(recall_score(labels, preds)))
45     print("Accuracy Score: {}".format(accuracy_score(labels, preds)))
46     print("F1 Score: {}".format(f1_score(labels, preds)))
47
48 # Find the optimal K value for KNN models.
49 def find_best_k(X_train, y_train, X_test, y_test, min_k=1, max_k=25):
50     best_k = 0
51     best_score = 0.0
52     for k in range(min_k, max_k+1, 2):
53         knn = KNeighborsClassifier(n_neighbors=k)
54         knn.fit(X_train, y_train)
55         preds = knn.predict(X_test)
56         f1 = f1_score(y_test, preds)

```

```

57         if f1 > best_score:
58             best_k = k
59             best_score = f1
60
61     print("Best Value for k: {}".format(best_k))
62     print("F1-Score: {}".format(best_score))
63
64     # Create a function that visualizes the confusion matrix for the model.
65     def plot_cm(model, normalize='true'):
66         fig, ax = plt.subplots(figsize=(8, 8))
67         plt.grid(False)
68         plot_confusion_matrix(model, X_test, y_test, cmap='Blues', ax=ax, normal
69
70     # Create function for performing log transformations.
71     def log_transform(df, features):
72         '''Runs a log transformation on a feature
73
74         @params
75         df is a pd.DataFrame
76         features is a list of columns to be considered
77
78         @output
79         new log-transformed column
80
81         '''
82     for feature in features:
83         df[feature + '_log'] = np.log(df[feature]+1)
84     return df

```

```

In [502]: 1 df = pd.read_csv('data/high_diamond_ranked_10min.csv')
          2
          3 df.columns

```

```

Out[502]: Index(['gameId', 'blueWins', 'blueWardsPlaced', 'blueWardsDestroyed',
                 'blueFirstBlood', 'blueKills', 'blueDeaths', 'blueAssists',
                 'blueEliteMonsters', 'blueDragons', 'blueHeralds',
                 'blueTowersDestroyed', 'blueTotalGold', 'blueAvgLevel',
                 'blueTotalExperience', 'blueTotalMinionsKilled',
                 'blueTotalJungleMinionsKilled', 'blueGoldDiff', 'blueExperienceDiff',
                 'blueCSPerMin', 'blueGoldPerMin', 'redWardsPlaced', 'redWardsDestroyed',
                 'redFirstBlood', 'redKills', 'redDeaths', 'redAssists',
                 'redEliteMonsters', 'redDragons', 'redHeralds', 'redTowersDestroyed',
                 'redTotalGold', 'redAvgLevel', 'redTotalExperience',
                 'redTotalMinionsKilled', 'redTotalJungleMinionsKilled', 'redGoldDiff',
                 'redExperienceDiff', 'redCSPerMin', 'redGoldPerMin'],
                 dtype='object')

```


In [503]:

1 df.head()

Out[503]:

	gameId	blueWins	blueWardsPlaced	blueWardsDestroyed	blueFirstBlood	blueKills	blueDea
0	4519157822	0	28	2	1	9	
1	4523371949	0	12	1	0	5	
2	4521474530	0	15	0	0	7	
3	4524384067	0	43	1	0	4	
4	4436033771	0	75	4	0	6	

5 rows × 40 columns



In [504]:

1 df.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 9879 entries, 0 to 9878

Data columns (total 40 columns):

#	Column	Non-Null Count	Dtype
0	gameId	9879 non-null	int64
1	blueWins	9879 non-null	int64
2	blueWardsPlaced	9879 non-null	int64
3	blueWardsDestroyed	9879 non-null	int64
4	blueFirstBlood	9879 non-null	int64
5	blueKills	9879 non-null	int64
6	blueDeaths	9879 non-null	int64
7	blueAssists	9879 non-null	int64
8	blueEliteMonsters	9879 non-null	int64
9	blueDragons	9879 non-null	int64
10	blueHeralds	9879 non-null	int64
11	blueTowersDestroyed	9879 non-null	int64
12	blueTotalGold	9879 non-null	int64
13	blueAvgLevel	9879 non-null	float64
14	blueTotalExperience	9879 non-null	int64
15	blueTotalMinionsKilled	9879 non-null	int64
16	blueTotalJungleMinionsKilled	9879 non-null	int64
17	blueGoldDiff	9879 non-null	int64
18	blueExperienceDiff	9879 non-null	int64
19	blueCSPerMin	9879 non-null	float64
20	blueGoldPerMin	9879 non-null	float64
21	redWardsPlaced	9879 non-null	int64
22	redWardsDestroyed	9879 non-null	int64
23	redFirstBlood	9879 non-null	int64
24	redKills	9879 non-null	int64
25	redDeaths	9879 non-null	int64
26	redAssists	9879 non-null	int64
27	redEliteMonsters	9879 non-null	int64
28	redDragons	9879 non-null	int64
29	redHeralds	9879 non-null	int64
30	redTowersDestroyed	9879 non-null	int64
31	redTotalGold	9879 non-null	int64
32	redAvgLevel	9879 non-null	float64
33	redTotalExperience	9879 non-null	int64
34	redTotalMinionsKilled	9879 non-null	int64
35	redTotalJungleMinionsKilled	9879 non-null	int64
36	redGoldDiff	9879 non-null	int64
37	redExperienceDiff	9879 non-null	int64
38	redCSPerMin	9879 non-null	float64
39	redGoldPerMin	9879 non-null	float64

dtypes: float64(6), int64(34)

memory usage: 3.0 MB


In [505]:

1 df.describe()

Out[505]:


	gameId	blueWins	blueWardsPlaced	blueWardsDestroyed	blueFirstBlood	blueK
count	9.879000e+03	9879.000000	9879.000000	9879.000000	9879.000000	9879.000000
mean	4.500084e+09	0.499038	22.288288	2.824881	0.504808	6.183900
std	2.757328e+07	0.500024	18.019177	2.174998	0.500002	3.011000
min	4.295358e+09	0.000000	5.000000	0.000000	0.000000	0.000000
25%	4.483301e+09	0.000000	14.000000	1.000000	0.000000	4.000000
50%	4.510920e+09	0.000000	16.000000	3.000000	1.000000	6.000000
75%	4.521733e+09	1.000000	20.000000	4.000000	1.000000	8.000000
max	4.527991e+09	1.000000	250.000000	27.000000	1.000000	22.000000

8 rows × 40 columns





```
In [506]: 1 fig = plt.figure(figsize=(4, 10))
          2 sns.heatmap(df.corr()[['blueWins']], annot=True)
```

Out[506]: <AxesSubplot:>



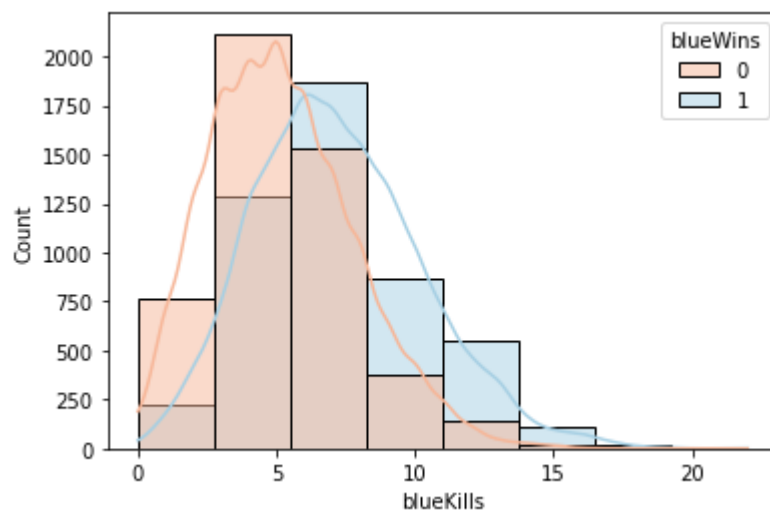
```
In [507]: 1 sns.histplot(x='blueGoldDiff', data=df, hue='blueWins', palette='RdBu', kde=
```

```
Out[507]: <AxesSubplot:xlabel='blueGoldDiff', ylabel='Count'>
```



```
In [508]: 1 sns.histplot(x='blueKills', data=df, hue='blueWins', palette='RdBu', kde=True
```

```
Out[508]: <AxesSubplot:xlabel='blueKills', ylabel='Count'>
```



```
In [ ]: 1
```

In []: 1

SCRUB

Data Preparation

Describe and justify the process for preparing the data for analysis.

Questions to consider:

- Were there variables you dropped or created?
- How did you address missing values or outliers?
- Why are these choices appropriate given the data and the business problem?

After initial data understanding, we are confident that the data we're using is sound. No NA or missing values were discovered.

Let's take a look at the columns:

In [509]: 1 df.columns

```
Out[509]: Index(['gameId', 'blueWins', 'blueWardsPlaced', 'blueWardsDestroyed',  
                'blueFirstBlood', 'blueKills', 'blueDeaths', 'blueAssists',  
                'blueEliteMonsters', 'blueDragons', 'blueHeralds',  
                'blueTowersDestroyed', 'blueTotalGold', 'blueAvgLevel',  
                'blueTotalExperience', 'blueTotalMinionsKilled',  
                'blueTotalJungleMinionsKilled', 'blueGoldDiff', 'blueExperienceDiff',  
                'blueCSPerMin', 'blueGoldPerMin', 'redWardsPlaced', 'redWardsDestroyed',  
                'redFirstBlood', 'redKills', 'redDeaths', 'redAssists',  
                'redEliteMonsters', 'redDragons', 'redHeralds', 'redTowersDestroyed',  
                'redTotalGold', 'redAvgLevel', 'redTotalExperience',  
                'redTotalMinionsKilled', 'redTotalJungleMinionsKilled', 'redGoldDiff',  
                'redExperienceDiff', 'redCSPerMin', 'redGoldPerMin'],  
              dtype='object')
```

There are a few columns that can be removed entirely and a few that can be combined into categorical variables.

First Blood

'First Blood' is awarded to the team who gets the first kill in the game. Both blueFirstBlood and redFirstBlood are binary and inversely related. If Blue wins First Blood, blueFirstBlood will be recorded as 1 and redFirstBlood will be recorded as 0.

We can merge these columns into one.

```
In [510]: 1 df['blueFirstBlood']
```

```
Out[510]: 0      1
          1      0
          2      0
          3      0
          4      0
          ..
        9874    1
        9875    0
        9876    0
        9877    1
        9878    1
        Name: blueFirstBlood, Length: 9879, dtype: int64
```

```
In [511]: 1 firstBlood = []
          2
          3 for item in df['blueFirstBlood']:
          4     if item == 1:
          5         firstBlood.append('Blue')
          6     else:
          7         firstBlood.append('Red')
          8
          9 df['firstBlood'] = firstBlood
         10
         11 df['firstBlood']
```

```
Out[511]: 0      Blue
          1      Red
          2      Red
          3      Red
          4      Red
          ...
        9874    Blue
        9875     Red
        9876     Red
        9877    Blue
        9878    Blue
        Name: firstBlood, Length: 9879, dtype: object
```

We can discard blueFirstBlood and redFirstBlood

```
In [512]: 1 del df['blueFirstBlood']
          2 del df['redFirstBlood']
```

Kills & Deaths

blueKills is inversely related with redDeaths, and redKills is inversely related with blueDeaths since the Blue team can only kill Red players and vice versa. blueDeaths and redDeaths can both be removed, leaving kills intact will preserve this information.

```
In [513]: 1 del df['blueDeaths']  
          2 del df['redDeaths']
```

Dragon & Herald

While this wouldn't hold true for LoL data spanning the entire length of each game, we know that there is only one opportunity to kill both the Dragon and the Herald in the first 10 minutes of each match. Unlike firstBlood where the action always occurs in the first 10 minutes (at least for the matches in our dataset), each dragon or herald can be killed only once or not at all.

Therefore, dragon and herald can be categorized as 'Blue,' 'Red,' or 'None.'

```
In [514]: 1 dragon_list = []  
          2  
          3 dragon_kill = df['blueDragons'] - df['redDragons']  
          4  
          5 for item in dragon_kill:  
          6     if item == 1:  
          7         dragon_list.append('Blue')  
          8     elif item == -1:  
          9         dragon_list.append('Red')  
         10     else:  
         11         dragon_list.append('No Dragon')  
         12  
         13 df['dragon'] = dragon_list
```

blueDragons and redDragons can be removed:

```
In [515]: 1 del df['blueDragons']  
          2 del df['redDragons']
```

We can reuse this code for the herald feature:

```
In [516]: 1 herald_list = []  
          2  
          3 herald_kill = df['blueHeralds'] - df['redHeralds']  
          4  
          5 for item in herald_kill:  
          6     if item == 1:  
          7         herald_list.append('Blue')  
          8     elif item == -1:  
          9         herald_list.append('Red')  
         10     else:  
         11         herald_list.append('No Herald')  
         12  
         13 df['herald'] = herald_list
```

```
In [517]: 1 del df['blueHeralds']  
          2 del df['redHeralds']
```

Elite Monsters

```
In [518]: 1 del df['blueEliteMonsters']
          2 del df['redEliteMonsters']
```

GoldDiff, ExperienceDiff, CSPerMin, and GoldPerMin

Both blue and red teams have these four metrics. While they are useful metrics for other types of analyses, they are essentially duplicative, since they are all calculated in a similar fashion from features already included in our data.

- GoldDiff represents the difference between blueTotalGold and redTotalGold
- ExperienceDiff represents the difference between blueTotalExperience and redTotalExperience
- blue and red CSPerMin represents the minute rate of blue and red TotalMinionsKilled. For our 10 minute data, CSPerMin for each team will always be TotalMinionsKilled divided by 10
- similarly, blue and red GoldPerMin represents blue and red TotalGold divided by 10

These four features from both teams (totaling 8 features) can be removed without losing any information.

```
In [519]: 1 del df['blueGoldDiff']
          2 del df['blueExperienceDiff']
          3 del df['blueCSPerMin']
          4 del df['blueGoldPerMin']
          5 del df['redGoldDiff']
          6 del df['redExperienceDiff']
          7 del df['redCSPerMin']
          8 del df['redGoldPerMin']
```

gameId

gameId represents a unique identifier for every LoL game, no two gameId's will ever be the same, so this column can be removed.

```
In [520]: 1 del df['gameId']
```

Reviewing cleaned data

In [521]:

```
1 print(df.columns)
2
3 print(df.shape)
```

```
Index(['blueWins', 'blueWardsPlaced', 'blueWardsDestroyed', 'blueKills',
      'blueAssists', 'blueTowersDestroyed', 'blueTotalGold', 'blueAvgLevel',
      'blueTotalExperience', 'blueTotalMinionsKilled',
      'blueTotalJungleMinionsKilled', 'redWardsPlaced', 'redWardsDestroyed',
      'redKills', 'redAssists', 'redTowersDestroyed', 'redTotalGold',
      'redAvgLevel', 'redTotalExperience', 'redTotalMinionsKilled',
      'redTotalJungleMinionsKilled', 'firstBlood', 'dragon', 'herald'],
      dtype='object')
(9879, 24)
```

We were able to remove 14 columns through this process without losing any information.

Alternative Dataset - Differences


In [522]:

```
1 df
```

Out[522]:

	blueWins	blueWardsPlaced	blueWardsDestroyed	blueKills	blueAssists	blueTowersDestroyed
0	0	28	2	9	11	(
1	0	12	1	5	5	(
2	0	15	0	7	4	(
3	0	43	1	4	5	(
4	0	75	4	6	6	(
...
9874	1	17	2	7	5	(
9875	1	54	0	6	8	(
9876	0	23	1	6	5	(
9877	0	14	4	2	3	(
9878	1	18	0	6	5	(

9879 rows × 24 columns



```

In [523]: 1 diff_df = pd.DataFrame()
          2
          3 diff_df['WardsPlaced'] = df['blueWardsPlaced'] - df['redWardsPlaced']
          4 diff_df['WardsDestroyed'] = df['blueWardsDestroyed'] - df['redWardsDestroyed']
          5 diff_df['Kills'] = df['blueKills'] - df['redKills']
          6 diff_df['Assists'] = df['blueAssists'] - df['redAssists']
          7 diff_df['TowersDestroyed'] = df['blueTowersDestroyed'] - df['redTowersDestroyed']
          8 diff_df['TotalGold'] = df['blueTotalGold'] - df['redTotalGold']
          9 diff_df['AvgLevel'] = df['blueAvgLevel'] - df['redAvgLevel']
         10 diff_df['TotalExperience'] = df['blueTotalExperience'] - df['redTotalExperience']
         11 diff_df['TotalMinionsKilled'] = df['blueTotalMinionsKilled'] - df['redTotalMinionsKilled']
         12 diff_df['TotalJungleMinionsKilled'] = df['blueTotalJungleMinionsKilled'] - df['redTotalJungleMinionsKilled']
         13
         14 diff_df = pd.concat([diff_df, df[['firstBlood', 'dragon', 'herald', 'blueWin']]])
         15
         16 diff_df.head()

```

Out[523]:

	WardsPlaced	WardsDestroyed	Kills	Assists	TowersDestroyed	TotalGold	AvgLevel	TotalExperience
0	13	-4	3	3	0	643	-0.2	
1	0	0	0	3	-1	-2908	-0.2	
2	0	-3	-4	-10	0	-1172	-0.4	
3	28	-1	-1	-5	0	-1321	0.0	
4	58	2	0	-1	0	-1004	0.0	

```

In [524]: 1 df = diff_df.copy()

```

```

In [ ]: 1

```

```

In [ ]: 1

```


EXPLORE


```

In [525]: 1 def heatmap(df_name, figsize=(15,10), cmap='Reds'):
2         corr = df_name.drop('blueWins',axis=1).corr()
3         mask = np.zeros_like(corr)
4         mask[np.triu_indices_from(mask)] = True
5         fig, ax = plt.subplots(figsize=figsize)
6         sns.heatmap(corr, annot=True, cmap=cmap, mask=mask)
7         return fig, ax
8
9 heatmap(df)

```

Out[525]: (<Figure size 1080x720 with 2 Axes>, <AxesSubplot:>)



```
In [526]: 1 # pd.set_option('display.max_rows', df.shape[0]+1)
2
3 # corr = df.drop('blueWins',axis=1).corr().abs()
4
5 # sort = corr.unstack()
6 # sort_order = sort.sort_values(kind="quicksort")
7
8 # corr_df = sort_order.to_frame()
9
10 # corr_features_df = corr_df[(corr_df[0] > 0.6) & (corr_df[0] < 1) ]
11
12 # corr_features_df.sort_values(by=0, ascending=False)
```

```
In [527]: 1 # https://pydatascience.org/2019/07/23/remove-duplicates-from-correlation-ma
2 def corr_list(df):
3     dataCorr = df.drop('blueWins',axis=1).corr()
4
5     dataCorr = dataCorr[abs(dataCorr) >= 0.01].stack().reset_index()
6     dataCorr = dataCorr[dataCorr['level_0'].astype(str)!=dataCorr['level_1']]
7
8     # filtering out lower/upper triangular duplicates
9     dataCorr['ordered-cols'] = dataCorr.apply(lambda x: '-'.join(sorted([x['
10     dataCorr = dataCorr.drop_duplicates(['ordered-cols']))
11     dataCorr.drop(['ordered-cols'], axis=1, inplace=True)
12
13     return dataCorr.sort_values(by=[0], ascending=False).head(10) #Get 10 hi
14
15 corr_list(df)
16
```

Out[527]:

	level_0	level_1	0
64	AvgLevel	TotalExperience	0.919161
23	Kills	TotalGold	0.917008
54	TotalGold	TotalExperience	0.894729
53	TotalGold	AvgLevel	0.833493
21	Kills	Assists	0.830751
25	Kills	TotalExperience	0.822845
24	Kills	AvgLevel	0.766222
33	Assists	TotalGold	0.759321
55	TotalGold	TotalMinionsKilled	0.638765
75	TotalExperience	TotalMinionsKilled	0.625556

Our multicollinearity analysis has presented a few variable relationships that need additional consideration.

- avgLevel and TotalExperience are highly correlated, which is not surprising. For now, we will stick with TotalExperience for blue and red teams since it's a bit more precise than AvgLevel.

However, we might want to run our baseline model with AvgLevel instead of TotalExperience to see if AvgLevel is more predictive.

- blue and red TotalGold appear consistently in our list. This is also not surprising since kills and assists award gold. We will remove gold for now.

In [528]:

```
1 corr_list(diff_df)
```

Out[528]:

	level_0	level_1	0
64	AvgLevel	TotalExperience	0.919161
23	Kills	TotalGold	0.917008
54	TotalGold	TotalExperience	0.894729
53	TotalGold	AvgLevel	0.833493
21	Kills	Assists	0.830751
25	Kills	TotalExperience	0.822845
24	Kills	AvgLevel	0.766222
33	Assists	TotalGold	0.759321
55	TotalGold	TotalMinionsKilled	0.638765
75	TotalExperience	TotalMinionsKilled	0.625556

```
In [529]: 1 df
```

Out[529]:

	WardsPlaced	WardsDestroyed	Kills	Assists	TowersDestroyed	TotalGold	AvgLevel	TotalEx
0	13	-4	3	3	0	643	-0.2	
1	0	0	0	3	-1	-2908	-0.2	
2	0	-3	-4	-10	0	-1172	-0.4	
3	28	-1	-1	-5	0	-1321	0.0	
4	58	2	0	-1	0	-1004	0.0	
...
9874	-29	-1	3	-2	0	2519	0.4	
9875	42	-21	2	5	0	782	0.2	
9876	9	1	-1	-6	0	-2416	-0.4	
9877	-52	0	-1	2	0	-839	-0.6	
9878	9	-2	0	1	0	927	0.2	

9879 rows × 14 columns

In [530]:

1df

Out[530]:

	WardsPlaced	WardsDestroyed	Kills	Assists	TowersDestroyed	TotalGold	AvgLevel	TotalEx
0	13	-4	3	3	0	643	-0.2	
1	0	0	0	3	-1	-2908	-0.2	
2	0	-3	-4	-10	0	-1172	-0.4	
3	28	-1	-1	-5	0	-1321	0.0	
4	58	2	0	-1	0	-1004	0.0	
...
9874	-29	-1	3	-2	0	2519	0.4	
9875	42	-21	2	5	0	782	0.2	
9876	9	1	-1	-6	0	-2416	-0.4	
9877	-52	0	-1	2	0	-839	-0.6	
9878	9	-2	0	1	0	927	0.2	

9879 rows × 14 columns



In []:

1

In []:

1

In []:

1

MODEL

In []:

1

Data Modeling

Describe and justify the process for analyzing or modeling the data.

Questions to consider:

- How did you analyze or model the data?
- How did you iterate on your initial approach to make it better?
- Why are these choices appropriate given the data and the business problem?

Train Test Split

```
In [531]: 1 def tt_split_df(df):
2
3     y = df['blueWins']
4     X = df.drop(columns=['blueWins'], axis=1)
5
6     X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
7
8     X_train_tf = X_train.copy()
9     X_test_tf = X_test.copy()
10
11     categoricals = X.select_dtypes('O').columns
12     numericals = X.select_dtypes('number').columns
13
14     encoder = OneHotEncoder(sparse=False, drop='if_binary')
15     train_categoricals = encoder.fit_transform(X_train_tf[categoricals])
16     test_categoricals = encoder.transform(X_test_tf[categoricals])
17
18     train_categoricals_df = pd.DataFrame(train_categoricals,
19                                         columns=encoder.get_feature_names(categ
20
21     test_categoricals_df = pd.DataFrame(test_categoricals,
22                                         columns=encoder.get_feature_names(categ
23     train_numericals_df = pd.DataFrame(scaler.fit_transform(X_train_tf[numer
24                                     columns=numericals))
25
26     test_numericals_df = pd.DataFrame(scaler.transform(X_test_tf[numericals]
27                                     columns=numericals))
28
29     X_train_tf = pd.concat([train_numericals_df, train_categoricals_df], axi
30     X_test_tf = pd.concat([test_numericals_df, test_categoricals_df], axis=1
31
32     return X_train_tf, X_test_tf, y_train, y_test
33
34 # tt_split_df(diff_df)
35
36 X_train, X_test, y_train, y_test = tt_split_df(diff_df)
37
38 X_train_tf, X_test_tf, y_train, y_test = tt_split_df(df)
39
```

```
In [532]: 1 X_train_tf
          2
```

Out[532]:

	WardsPlaced	WardsDestroyed	Kills	Assists	TowersDestroyed	TotalGold	AvgLevel
0	-0.791231	-0.380079	-0.245136	-0.337819	-0.019208	-0.163406	-0.398722
1	0.003580	1.043867	0.471110	0.183071	-0.019208	0.306376	1.698298
2	0.154972	-0.024093	0.709858	-0.337819	3.069716	1.276605	1.698298
3	-0.980472	-1.804026	-0.961381	-0.858708	-0.019208	-0.876870	0.020682
4	-0.034269	-1.448039	-0.961381	-0.511449	-0.019208	-0.721911	-0.818126
...
6910	0.003580	-0.736066	-0.722633	-1.032338	-0.019208	-1.157349	-0.818126
6911	0.041428	0.331894	0.471110	0.356701	-0.019208	0.727912	0.859490
6912	-0.450598	1.043867	-1.200130	-1.205968	-0.019208	-1.100926	-0.818126
6913	-0.450598	-0.024093	-1.916375	-0.858708	-0.019208	-1.756331	-2.495742
6914	-0.677687	-0.736066	0.471110	0.530331	-0.019208	0.149782	0.020682

6915 rows × 17 columns

```
In [ ]: 1
```

```
In [ ]: 1
```

```
In [ ]: 1
```

```
In [ ]: 1
```

```
In [ ]: 1
```

```
In [ ]: 1
```

```
In [ ]: 1
```

```
In [533]: 1 # >>> X_train, X_test, _, _ = train_test_split(X,y, test_size=0.33, random_s
```

```
In [534]: 1 # Identify features and target
2
3 y = df['blueWins']
4 X = df.drop(columns=['blueWins'], axis=1)
5
6 # Assign train / test split
7
8 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, ran
9
10 # Confirm split
11
12 print("X_train shape:", X_train.shape)
13 print("X_test shape:", X_test.shape)
```

X_train shape: (6915, 13)

X_test shape: (2964, 13)

```
In [ ]: 1
```

```
In [ ]: 1
```

```
In [ ]: 1
```

Transforming X Train and Test

```
In [535]: 1 # X_train_tf = X_train.copy()
2 # X_test_tf = X_test.copy()
```

Categorical Columns

```
In [536]: 1 categoricals = X.select_dtypes('O').columns
2 numericals = X.select_dtypes('number').columns
3 categoricals, numericals
```

```
Out[536]: (Index(['firstBlood', 'dragon', 'herald'], dtype='object'),
Index(['WardsPlaced', 'WardsDestroyed', 'Kills', 'Assists', 'TowersDestroyed',
'TotalGold', 'AvgLevel', 'TotalExperience', 'TotalMinionsKilled',
'TotalJungleMinionsKilled'],
dtype='object'))
```



```
In [537]: 1  ## Encode categorical columns, only drop if binary
2  encoder = OneHotEncoder(sparse=False, drop='if_binary')
3  train_categoricals = encoder.fit_transform(X_train[categoricals])
4  test_categoricals = encoder.transform(X_test[categoricals])
5  train_categoricals
```

```
Out[537]: array([[0., 0., 1., ..., 0., 1., 0.],
 [0., 0., 0., ..., 0., 1., 0.],
 [1., 0., 0., ..., 1., 0., 0.],
 ...,
 [1., 0., 0., ..., 0., 1., 0.],
 [1., 1., 0., ..., 0., 1., 0.],
 [1., 0., 0., ..., 1., 0., 0.]])
```

```
In [538]: 1  train_categoricals_df = pd.DataFrame(train_categoricals,
2                                             columns=encoder.get_feature_names(categ
3
4  test_categoricals_df = pd.DataFrame(test_categoricals,
5                                       columns=encoder.get_feature_names(categ
6
7  train_categoricals_df.head()
```

```
Out[538]:
```

	firstBlood_Red	dragon_Blue	dragon_No Dragon	dragon_Red	herald_Blue	herald_No Herald	herald_Red
0	0.0	0.0	1.0	0.0	0.0	1.0	0.0
1	0.0	0.0	0.0	1.0	0.0	1.0	0.0
2	1.0	0.0	0.0	1.0	1.0	0.0	0.0
3	0.0	0.0	0.0	1.0	0.0	0.0	1.0
4	1.0	1.0	0.0	0.0	0.0	1.0	0.0

Numerical Columns

```
In [539]: 1 from sklearn.preprocessing import StandardScaler
2
3 scaler = scaler = StandardScaler()
4
5 train_numericals_df = pd.DataFrame(scaler.fit_transform(X_train[numericals])
6                                   columns=numericals)
7
8 test_numericals_df = pd.DataFrame(scaler.transform(X_test[numericals]),
9                                   columns=numericals)
10
11 train_numericals_df.head()
```

Out[539]:

	WardsPlaced	WardsDestroyed	Kills	Assists	TowersDestroyed	TotalGold	AvgLevel	TotalKills
0	-0.312890	1.373334	0.220704	-0.349942	-0.04164	-0.001092	-0.821423	0.000000
1	0.156072	-0.741644	0.220704	-0.003060	-0.04164	-0.440888	-0.821423	0.000000
2	-0.000249	-0.389148	-0.733301	-0.349942	-0.04164	-0.276167	-0.405661	0.000000
3	0.156072	1.373334	2.128716	2.425111	-0.04164	2.063192	1.673149	0.000000
4	-0.039329	-0.036651	0.697707	0.343822	-0.04164	0.381093	0.010101	0.000000

```
In [540]: 1 X_train = pd.concat([train_numericals_df, train_categoricals_df], axis=1)
2 X_test = pd.concat([test_numericals_df, test_categoricals_df], axis=1)
3
4 X_train.head()
```

Out[540]:

	WardsPlaced	WardsDestroyed	Kills	Assists	TowersDestroyed	TotalGold	AvgLevel	TotalKills
0	-0.312890	1.373334	0.220704	-0.349942	-0.04164	-0.001092	-0.821423	0.000000
1	0.156072	-0.741644	0.220704	-0.003060	-0.04164	-0.440888	-0.821423	0.000000
2	-0.000249	-0.389148	-0.733301	-0.349942	-0.04164	-0.276167	-0.405661	0.000000
3	0.156072	1.373334	2.128716	2.425111	-0.04164	2.063192	1.673149	0.000000
4	-0.039329	-0.036651	0.697707	0.343822	-0.04164	0.381093	0.010101	0.000000

Logistic Regression

```
In [541]: 1 model_log = LogisticRegression()
          2
          3 model_log.fit(X_train, y_train)
          4
          5 # scores = cross_val_score(model_lr, X_train,y_train, cv=10) # model, train,
          6
          7
```

Out[541]: LogisticRegression()

```
In [542]: 1 y_train.value_counts(1)
```

Out[542]: 1 0.500217
0 0.499783
Name: blueWins, dtype: float64

```
In [543]: 1 y_test.value_counts(1)
```

Out[543]: 0 0.503711
1 0.496289
Name: blueWins, dtype: float64

```
In [544]: 1 ## Get Predictions for training and test data to check metrics functions
          2 y_hat_train = model_log.predict(X_train)
          3 y_hat_test = model_log.predict(X_test)
```

Model Accuracy

```
In [545]: 1 def model_accuracy(model_string_name, model, X_train=X_train, y_train=y_train)
          2     print(f'{model_string_name} accuracy score:')
          3     print(f'Training Score:\t{model_log.score(X_train,y_train):.2f}')
          4     print(f'Test Score:\t{model_log.score(X_test,y_test):.2f}')
          5
          6 model_accuracy('Logistic Regression', model_log)
```

Logistic Regression accuracy score:
Training Score: 0.74
Test Score: 0.72

Cross Validation Check

```
In [546]: 1 def cross_val_check(model_string_name, model, X_train=X_train, y_train=y_train, cv=10):
2         scores = cross_val_score(model, X_train, y_train, cv=10) # model, train,
3         print(f'{model_string_name} Cross Validation Scores:\n')
4         print(scores)
5         print(f'\nCross validation mean: \t{scores.mean():.3f}')
6
7         cross_val_check('Logistic Regression', model_log)
```

Logistic Regression Cross Validation Scores:

```
[0.72976879 0.73265896 0.75867052 0.73410405 0.73988439 0.74384949
 0.72503618 0.75253256 0.72648336 0.73950796]
```

Cross validation mean: 0.738

Confusion Matrix

In []:

1

In [547]:

```

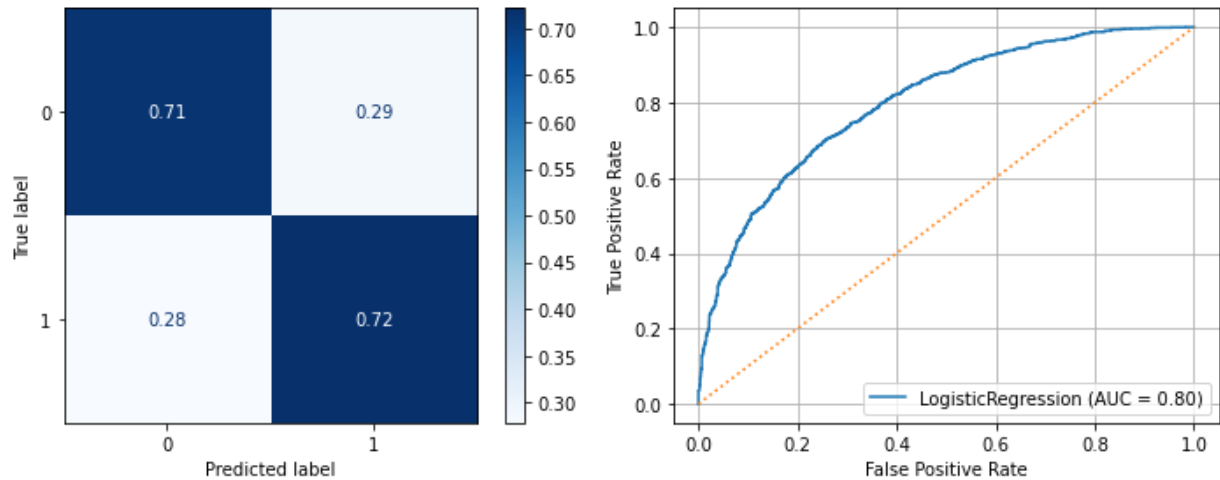
1  ## Modified version of our simple eval function from Topic 25 Part 2 Study G
2  # - Added X_train and y_train for if we want scores for both train and test
3  def evaluate_classification(model, X_test_tf,y_test,cmap='Blues',
4                             normalize='true',classes=None,figsize=(10,4),
5                             X_train = None, y_train = None,):
6      """Evaluates a scikit-learn binary classification model.
7
8      Args:
9          model ([type]): [description]
10         X_test_tf ([type]): [description]
11         y_test ([type]): [description]
12         cmap (str, optional): [description]. Defaults to 'Reds'.
13         normalize (str, optional): [description]. Defaults to 'true'.
14         classes ([type], optional): [description]. Defaults to None.
15         figsize (tuple, optional): [description]. Defaults to (8,4).
16         X_train ([type], optional): [description]. Defaults to None.
17         y_train ([type], optional): [description]. Defaults to None.
18     """
19
20
21     y_hat_test = model.predict(X_test_tf)
22     print(metrics.classification_report(y_test, y_hat_test,target_names=clas
23
24     fig,ax = plt.subplots(ncols=2,figsize=figsize)
25     metrics.plot_confusion_matrix(model, X_test_tf,y_test,cmap=cmap,
26                                  normalize=normalize,display_labels=classes
27                                  ax=ax[0])
28
29     curve = metrics.plot_roc_curve(model,X_test_tf,y_test,ax=ax[1])
30     curve.ax_.grid()
31     curve.ax_.plot([0,1],[0,1],ls=':')
32     fig.tight_layout()
33     plt.show()
34
35     ## Add comparing Scores if X_train and y_train provided.
36     if (X_train is not None) & (y_train is not None):
37         print(f"Training Score = {model.score(X_train,y_train):.2f}")
38         print(f"Test Score = {model.score(X_test_tf,y_test):.2f}")
39
40
41
42 def evaluate_grid(grid,X_test,y_test,X_train=None,y_train=None):
43     print('The best parameters were:')
44     print("\t",grid.best_params_)
45
46     model = grid.best_estimator_
47
48     print('\n[i] Classification Report')
49     evaluate_classification(model, X_test,y_test,X_train=X_train,y_train=y_t

```

In [548]:

1 evaluate_classification(model_log, X_test, y_test)

	precision	recall	f1-score	support
0	0.72	0.71	0.72	1493
1	0.71	0.72	0.72	1471
accuracy			0.72	2964
macro avg	0.72	0.72	0.72	2964
weighted avg	0.72	0.72	0.72	2964



Dummy Check